

Given historical data on loans given out with information on whether or not the borrower defaulted (charge-off), can we build a model that can predict whether or not a borrower will pay back their loan? This way in the future when we get a new potential customer we can assess whether or not they are likely to pay back the loan. Keep in mind classification metrics when evaluating the performance of your model! source: <https://www.kaggle.com/harlfoxem/housesalesprediction>

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

In [2]:

```
Data=pd.read_csv(r'C:\Users\chumj\Downloads\lending.csv')
```

In [3]:

```
Data.head(3)
```

Out[3]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ownership	annual_inc	...	open_acc	pu
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	RENT	117000.0	...	16.0	
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	MORTGAGE	65000.0	...	17.0	
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	RENT	43057.0	...	13.0	

3 rows × 27 columns

In [4]:

```
Data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            396030 non-null float64
1   term                                396030 non-null object
2   int_rate                             396030 non-null float64
3   installment                           396030 non-null float64
4   grade                                396030 non-null object
5   sub_grade                            396030 non-null object
6   emp_title                            373103 non-null object
7   emp_length                           377729 non-null object
8   home_ownership                       396030 non-null object
9   annual_inc                           396030 non-null float64
10  verification_status                  396030 non-null object
11  issue_d                              396030 non-null object
12  loan_status                           396030 non-null object
13  purpose                              396030 non-null object
14  title                                394275 non-null object
15  dti                                   396030 non-null float64
16  earliest_cr_line                      396030 non-null object
17  open_acc                              396030 non-null float64
18  pub_rec                               396030 non-null float64
19  revol_bal                             396030 non-null float64
20  total_acc                             396030 non-null float64
21  delinq_12m                            396030 non-null float64
22  late_12m                              396030 non-null float64
23  charge_off                            396030 non-null float64
24  charge_off_ratio                      396030 non-null float64
25  dti_max                               396030 non-null float64
26  dti_min                               396030 non-null float64
27  dti_max_ratio                         396030 non-null float64
```

```

20  revol_util          395754 non-null  float64
21  total_acc           396030 non-null  float64
22  initial_list_status  396030 non-null  object
23  application_type     396030 non-null  object
24  mort_acc            358235 non-null  float64
25  pub_rec_bankruptcies  395495 non-null  float64
26  address             396030 non-null  object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB

```

## EXPLORATORY DATA ANALYSIS.

Keep in mind our soul aim is to predict if a customer will predict or default a loan. Let's start by exploring and anylsing some import faetures in our Dataset.The loan status itself is a very important feature,lets start with it.

In [5]:

```
Data.columns
```

Out[5]:


```

Index(['loan_amnt', 'term', 'int_rate', 'installment', 'grade', 'sub_grade',
      'emp_title', 'emp_length', 'home_ownership', 'annual_inc',
      'verification_status', 'issue_d', 'loan_status', 'purpose', 'title',
      'dti', 'earliest_cr_line', 'open_acc', 'pub_rec', 'revol_bal',
      'revol_util', 'total_acc', 'initial_list_status', 'application_type',
      'mort_acc', 'pub_rec_bankruptcies', 'address'],
      dtype='object')

```

In [6]:

```
sns.countplot(x='loan_status',data=Data);
```




In [7]:

```

#loan amount
sns.distplot(Data['loan_amnt'],bins=50,kde=False);

```



In [8]:

```
#Launching correlation between the continous variable features.  
Data.corr()
```

Out[8]:

	loan_amnt	int_rate	installment	annual_inc	dti	open_acc	pub_rec	revol_bal	revol_util	total_acc	mo
loan_amnt	1.000000	0.168921	0.953929	0.336887	0.016636	0.198556	0.077779	0.328320	0.099911	0.223886	0.2
int_rate	0.168921	1.000000	0.162758	-0.056771	0.079038	0.011649	0.060986	0.011280	0.293659	0.036404	0.0
installment	0.953929	0.162758	1.000000	0.330381	0.015786	0.188973	0.067892	0.316455	0.123915	0.202430	0.1
annual_inc	0.336887	0.056771	0.330381	1.000000	0.081685	0.136150	0.013720	0.299773	0.027871	0.193023	0.2
dti	0.016636	0.079038	0.015786	-0.081685	1.000000	0.136181	0.017639	0.063571	0.088375	0.102128	0.0
open_acc	0.198556	0.011649	0.188973	0.136150	0.136181	1.000000	0.018392	0.221192	0.131420	0.680728	0.1
pub_rec	-0.077779	0.060986	-0.067892	-0.013720	0.017639	-0.018392	1.000000	0.101664	0.075910	0.019723	0.0
revol_bal	0.328320	0.011280	0.316455	0.299773	0.063571	0.221192	0.101664	1.000000	0.226346	0.191616	0.1
revol_util	0.099911	0.293659	0.123915	0.027871	0.088375	-0.131420	0.075910	0.226346	1.000000	0.104273	0.0
total_acc	0.223886	0.036404	0.202430	0.193023	0.102128	0.680728	0.019723	0.191616	0.104273	1.000000	0.3
mort_acc	0.222315	0.082583	0.193694	0.236320	0.025439	0.109205	0.011552	0.194925	0.007514	0.381072	1.0
pub_rec_bankruptcies	-0.106539	0.057450	-0.098628	-0.050162	0.014558	-0.027732	0.699408	0.124532	0.086751	0.042035	0.0

In [9]:


```
Data.corr()['loan_amnt'].sort_values(ascending=False)
```


Out[9]:

```
loan_amnt      1.000000  
installment    0.953929  
annual_inc     0.336887  
revol_bal      0.328320  
total_acc      0.223886  
mort_acc       0.222315  
open_acc       0.198556  
int_rate       0.168921  
revol_util     0.099911  
dti            0.016636  
pub_rec       -0.077779  
pub_rec_bankruptcies -0.106539  
Name: loan_amnt, dtype: float64
```

In [10]:


```
plt.figure(figsize=(10,8))  
sns.heatmap(Data.corr(),annot=True,cmap='viridis');
```





In [11]:

```
#installment and loan amout seems to correlation well,lets take a look for more Analysis
sns.scatterplot(x='installment',y='loan_amnt',data=Data);
```




lets take look at loan sttus(categorical feature) and loan amount.

In [12]:

```
sns.barplot(x='loan_status',y='loan_amnt',data=Data)
```

Out[12]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x18ce36ad4c8>




In [13]:

```
sns.boxplot(x='loan_status',y='loan_amnt',data=Data)
```

Out[13]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x18ce370c9c8>



In [14]:

```
Data.groupby('loan_status')['loan_amnt'].describe()
```

Out[14]:

	count	mean	std	min	25%	50%	75%	max
loan_status								
Charged Off	77673.0	15126.300967	8505.090557	1000.0	8525.0	14000.0	20000.0	40000.0
Fully Paid	318357.0	13866.878771	8302.319699	500.0	7500.0	12000.0	19225.0	40000.0

In [15]:

```
#sns.kdeplot(Data['installment'],Data['loan_amnt'],  
             #cmap="plasma", shade=True, shade_lowest=False);
```

In [ ]:

In [16]:

```
#lets take a look of the Grading and subGrading in relation to the loan
```

In [17]:

```
sorted(Data['grade'].unique())
```

Out[17]:

```
['A', 'B', 'C', 'D', 'E', 'F', 'G']
```

In [18]:

```
sorted(Data['sub_grade'].unique())
```

Out[18]:

```
['A1',  
 'A2',  
 'A3']
```

```


'A3',
'A4',
'A5',
'B1',
'B2',
'B3',
'B4',
'B5',
'C1',
'C2',
'C3',
'C4',
'C5',
'D1',
'D2',
'D3',
'D4',
'D5',
'E1',
'E2',
'E3',
'E4',
'E5',
'F1',
'F2',
'F3',
'F4',
'F5',
'G1',
'G2',
'G3',
'G4',
'G5']

```

let's check the grading with relation to loan status graphically

In [19]:

```
sns.countplot(x='grade',hue='loan_status',data=Data);
```





In [20]:

```

plt.figure(figsize=(10,8))
subgrade_order=sorted(Data['sub_grade'].unique())
sns.countplot(x='sub_grade',data=Data,order=subgrade_order,hue='loan_status');

```





something seems to be happening with F and G, It seems they do not paid back that often

In [21]:

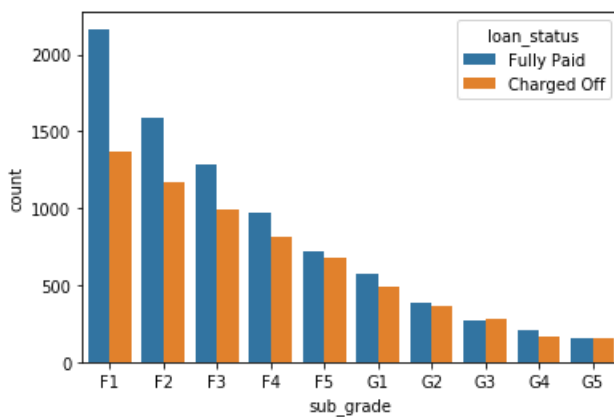
```
F_and_G = Data[(Data['grade'] == 'G') | (Data['grade'] == 'F')]
subgrade_order = sorted(F_and_G['sub_grade'].unique())
```

In [22]:

```
sns.countplot(x='sub_grade', order=subgrade_order, hue='loan_status', data=F_and_G)
```

Out[22]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x18ce16e2648>



Loan status recall is the main feature we are about to predict, let's change "Fully Paid"=1, "Charged Off"=0

In [23]:

```
Data['loan_status'].unique()
```

Out[23]:

```
array(['Fully Paid', 'Charged Off'], dtype=object)
```

In [24]:

```
Data['repaid_loans'] = Data['loan_status'].map({'Fully Paid': 1, 'Charged Off': 0})
```

In [25]:

```
Data[['repaid_loans', 'loan_status']].head(5)
```

Out[25]:

	repaid_loans	loan_status
0	1	Fully Paid
1	1	Fully Paid
2	1	Fully Paid
3	1	Fully Paid
4	0	Charged Off

In [26]:


```
Data['repaid_loans']
```

Out[26]:

```
0      1
1      1
2      1
3      1
4      0
..
396025  1
396026  1
396027  1
396028  1
396029  1
Name: repaid_loans, Length: 396030, dtype: int64
```

In [27]:

```
Data.corr()['repaid_loans'].sort_values().drop('repaid_loans').plot(kind='bar');
```



In [28]:

```
Data.corr()['repaid_loans'].sort_values().drop('repaid_loans')
```

Out[28]:

```
int_rate      -0.247758
revol_util    -0.082373
dti           -0.062413
loan_amnt     -0.059836
installment   -0.041082
open_acc      -0.028012
pub_rec       -0.019933
pub_rec_bankruptcies -0.009383
revol_bal      0.010892
total_acc      0.017893
annual_inc     0.053432
mort_acc       0.062111
```



mort\_acc 0.0/3111  
Name: repaid\_loans, dtype: float64

# PREPROCESSING OF OUR DATASET

Converting categorical strings features to dummy variable,removing and filling missing data

In [29]:

```
Data.head(2)
```

Out[29]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ownership	annual_inc	...	pub_rec	revo
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	RENT	117000.0	...	0.0	36%
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	MORTGAGE	65000.0	...	0.0	20%

2 rows × 28 columns

In [30]:

```
len(Data)
```

Out[30]:

396030

In [31]:

```
Data.isnull().sum()
```

Out[31]:


```
loan_amnt      0
term            0
int_rate       0
installment    0
grade          0
sub_grade      0
emp_title      22927
emp_length     18301
home_ownership  0
annual_inc     0
verification_status  0
issue_d        0
loan_status    0
purpose        0
title          1755
dti            0
earliest_cr_line  0
open_acc       0
pub_rec        0
revol_bal      0
revol_util     276
total_acc      0
initial_list_status  0
application_type  0
mort_acc       37795
pub_rec_bankruptcies  535
address        0
repaid_loans   0
dtype: int64
```

In [32]:

```
plt.figure(figsize=(10,8))
sns.heatmap(Data.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

Out[32]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x18ce1227f88>



In [33]:

```
# percentage of the entire dataset
Data.isnull().sum()/len(Data)*100
```

Out[33]:

loan_amnt	0.000000
term	0.000000
int_rate	0.000000
installment	0.000000
grade	0.000000
sub_grade	0.000000
emp_title	5.789208
emp_length	4.621115
home_ownership	0.000000
annual_inc	0.000000
verification_status	0.000000
issue_d	0.000000
loan_status	0.000000
purpose	0.000000
title	0.443148
dti	0.000000
earliest_cr_line	0.000000
open_acc	0.000000
pub_rec	0.000000
revol_bal	0.000000
revol_util	0.069692
total_acc	0.000000
initial_list_status	0.000000
application_type	0.000000
mort_acc	9.543469
pub_rec_bankruptcies	0.135091
address	0.000000

```
repaid_loans          0.000000
dtype: float64
```

In [34]:

```
# let's start with mort_acc of almost 10%,we cannot drop it out
Data['mort_acc'].value_counts()
```

Out[34]:

```
0.0      139777
1.0       60416
2.0       49948
3.0       38049
4.0       27887
5.0       18194
6.0       11069
7.0        6052
8.0        3121
9.0        1656
10.0         865
11.0         479
12.0         264
13.0         146
14.0         107
15.0          61
16.0          37
17.0          22
18.0          18
19.0          15
20.0          13
24.0          10
22.0           7
21.0           4
25.0           4
27.0           3
23.0           2
32.0           2
26.0           2
31.0           2
30.0           1
28.0           1
34.0           1
Name: mort_acc, dtype: int64
```

In [35]:


```
#looking which feature correlate most with mort_acc,then select teh faecture and look for it's mean
#and later apply it to the missing values in mort_acc
Data.corr()['mort_acc'].sort_values()
```

Out[35]:

```
int_rate          -0.082583
dti               -0.025439
revol_util         0.007514
pub_rec           0.011552
pub_rec_bankruptcies 0.027239
repaid_loans       0.073111
open_acc          0.109205
installment       0.193694
revol_bal         0.194925
loan_amnt         0.222315
annual_inc        0.236320
total_acc         0.381072
mort_acc          1.000000
Name: mort_acc, dtype: float64
```

In [36]:

```
Data.corr()['mort_acc'].sort_values().drop('mort_acc').plot(kind='bar');
```



In [37]:

```
#total_acc feautres correlate most with mort_acc,let use the fillna() by groupby the total_acc
#and calculate mean value for the mort_acc per entry.
print('mean of mort_acc col per total_acc')
Data.groupby('total_acc').mean()['mort_acc']
```

mean of mort\_acc col per total\_acc

Out[37]:

```
total_acc
2.0      0.000000
3.0      0.052023
4.0      0.066743
5.0      0.103289
6.0      0.151293
...
124.0    1.000000
129.0    1.000000
135.0    3.000000
150.0    2.000000
151.0    0.000000
Name: mort_acc, Length: 118, dtype: float64
```

In [38]:

```
# if mort_acc is missing,we fill in that missing value with the mean value corresponding to it's total_acc
total_acc_avg=Data.groupby('total_acc').mean()['mort_acc']
```

In [39]:

```
total_acc_avg[2]
# we can choose,2=0,3=0.052023 etc
```

Out[39]:

0.0

In [40]:

```
def fill_mort_acc(total_acc,mort_acc):
    if np.isnan(mort_acc):
        return total_acc_avg[total_acc]
    else:
        return mort_acc
```

In [41]:

```
Data['mort_acc']=Data.apply(lambda x:fill_mort_acc(x['total_acc'],x['mort_acc']),axis=1)
```

In [42]:

```
Data.isnull().sum()
```

Out[42]:

```
loan_amnt      0
term           0
int_rate       0
installment    0
grade          0
sub_grade      0
emp_title      22927
emp_length     18301
home_ownership 0
annual_inc     0
verification_status 0
issue_d        0
loan_status    0
purpose        0
title          1755
dti            0
earliest_cr_line 0
open_acc       0
pub_rec        0
revol_bal      0
revol_util     276
total_acc      0
initial_list_status 0
application_type 0
mort_acc       0
pub_rec_bankruptcies 535
address        0
repaid_loans   0
dtype: int64
```

In [43]:

```
#let's take a look at emp_title
Data['emp_title'].nunique()
```

Out[43]:

```
173105
```

In [44]:

```
Data['emp_title'].value_counts()
#from below it shows many emp_title and is not realistic to convert them to a dummy variable,it makes sense if we drop it.
```

Out[44]:

```
Teacher      4389
Manager      4250
Registered Nurse 1856
RN           1846
Supervisor   1830
...
Dore Academy      1
Loan servicing    1
Director of Quality And Admin 1
admitting registrar 1
Rockford Police Department 1
Name: emp_title, Length: 173105, dtype: int64
```

In [45]:

```
Data=Data.drop('emp_title',axis=1)
```

In [46]:

```
Data['emp_length'].value_counts()
```

Out[46]:

```
10+ years    126041
2 years      35827
< 1 year     31725
3 years      31665
5 years      26495
1 year       25882
4 years      23952
6 years      20841
7 years      20819
8 years      19168
9 years      15314
Name: emp_length, dtype: int64
```

In [47]:

```
sorted(Data['emp_length'].dropna().unique())
```

Out[47]:


```
['1 year',
 '10+ years',
 '2 years',
 '3 years',
 '4 years',
 '5 years',
 '6 years',
 '7 years',
 '8 years',
 '9 years',
 '< 1 year']
```


In [48]:

```
emp_length_order= ['< 1 year','1 year','2 years', '3 years', '4 years','5 years','6 years','7 years','8 years','9 years','10+ years',]
```

In [49]:


```
plt.figure(figsize=(10,8))
sns.countplot(x='emp_length',order=emp_length_order,data=Data);
```





In [50]:

```
plt.figure(figsize=(10,3))
sns.countplot(x='emp_length',order=emp_length_order,data=Data,hue='loan_status');
# looking at the plot below,charged off seems to be even from <1year....10+years,no real pattern w
e can exactract from emp_length,
#so it makes sense if we drop it.
```



In [51]:

```
Data=Data.drop('emp_length',axis=1)
```

In [52]:

```
# checking title
Data['title'].value_counts()
```

Out[52]:

```
Debt consolidation      152472
Credit card refinancing  51487
Home improvement        15264
Other                   12930
Debt Consolidation      11608
...
Moving Cost              1
To North Dakota          1
PAY LESS                 1
2013 Life Change         1
Save my sanity           1
Name: title, Length: 48817, dtype: int64
```

In [53]:

```
Data['title'].unique()
```

Out[53]:

```
array(['Vacation', 'Debt consolidation', 'Credit card refinancing', ...,
      'Credit buster ', 'Loanforpayoff', 'Toxic Debt Payoff'],
      dtype=object)
```

In [54]:

```
Data['purpose']
```

Out[54]:

```
0          vacation
1    debt_consolidation
2          credit_card
3          credit_card
4          credit_card
```

```
4          credit_card
      ...
396025    debt_consolidation
396026    debt_consolidation
396027    debt_consolidation
396028    debt_consolidation
396029    debt_consolidation
Name: purpose, Length: 396030, dtype: object
```

Title and purpose seems to be the same information, dropping title makes sense.

In [55]:

```
Data = Data.drop('title', axis=1)
```

In [56]:

```
Data.isnull().sum()
```

Out[56]:

```
loan_amnt      0
term           0
int_rate       0
installment    0
grade          0
sub_grade      0
home_ownership 0
annual_inc     0
verification_status 0
issue_d        0
loan_status    0
purpose        0
dti            0
earliest_cr_line 0
open_acc       0
pub_rec        0
revol_bal      0
revol_util     276
total_acc      0
initial_list_status 0
application_type 0
mort_acc       0
pub_rec_bankruptcies 535
address        0
repaid_loans    0
dtype: int64
```

revol\_util and pub\_rec\_bankruptcies just record less than 0.5% so it insignificant, so we can drop them.

In [57]:

```
Data = Data.dropna()
```

In [58]:

```
#Data.isnull().sum()
```

In [59]:

```
Data.select_dtypes(['object']).columns
```

Out[59]:

```
Index(['term', 'grade', 'sub_grade', 'home_ownership', 'verification_status',
      'issue_d', 'loan_status', 'purpose', 'earliest_cr_line',
      'initial_list_status', 'application_type', 'address'],
      dtype='object')
```



## let's take a look of all the string features

In [60]:

```
Data['term'].nunique()
```

Out[60]:

2

In [61]:

```
Data['term'].value_counts()
```

Out[61]:

```
36 months    301247
60 months     93972
Name: term, dtype: int64
```

In [62]:

```
Data['term']=Data['term'].apply(lambda term:int(term[:3]))
```

In [63]:

```
#grade is a sub_grade,so we can drop the grade feature.
Data=Data.drop('grade',axis=1)
```

In [64]:

```
Data['home_ownership'].nunique()
```

Out[64]:

6

In [65]:

```
Data['home_ownership'].value_counts()
```

Out[65]:

```
MORTGAGE    198022
RENT         159395
OWN          37660
OTHER         110
NONE          29
ANY           3
Name: home_ownership, dtype: int64
```

In [66]:

```
# we can convert home_ownership to dummy variable,lets try to combine 'none 'and 'any' to 'Other'
```

In [67]:

```
Data['home_ownership']=Data['home_ownership'].replace(['NONE','ANY'],'OTHER')
Dummies=pd.get_dummies(Data['home_ownership'],drop_first=True)
Data=Data.drop('home_ownership',axis=1)
Data=pd.concat([Data,Dummies],axis=1)
```

In [68]:

```
#for the address col we can extract the zip code from the address,we can make a new col call zip
Data['zip']=Data['address'].apply(lambda address:address[-5:])
```

In [69]:

```
Data['address'].head(2)
```

Out[69]:

```
0      0174 Michelle Gateway\r\nMendozaberg, OK 22690
1      1076 Carney Fort Apt. 347\r\nLoganmouth, SD 05113
Name: address, dtype: object
```

In [70]:

```
Dummies=pd.get_dummies(Data['zip'],drop_first=True)
Data=Data.drop(['zip','address'],axis=1)
Data=pd.concat([Data,Dummies],axis=1)
```

In [71]:

```
Data['issue_d'].head(3)
#this can be track or leakage,telling us ahead wheather loan was issue or not,so we immediately dr
op it out
```

Out[71]:

```
0      Jan-2015
1      Jan-2015
2      Jan-2015
Name: issue_d, dtype: object
```

In [72]:

```
Data=Data.drop('issue_d',axis=1)
```

In [73]:

```
#This like a historic time stamp which we can apply lambda function for the years
Data['earliest_year']=Data['earliest_cr_line'].apply(lambda date:int(date[-4]))
```

In [74]:

```
Data['earliest_cr_line'].head(2)
```

Out[74]:

```
0      Jun-1990
1      Jul-2004
Name: earliest_cr_line, dtype: object
```

In [75]:

```
Data=Data.drop('earliest_cr_line',axis=1)
```

In [76]:

```
Data.head(2)
```

Out[76]:

	loan_amnt	term	int_rate	installment	sub_grade	annual_inc	verification_status	loan_status	purpose	dti	...	05113
0	10000.0	36	11.44	329.48	B4	117000.0	Not Verified	Fully Paid	vacation	26.24	...	0
1	8000.0	36	11.99	265.68	B5	65000.0	Not Verified	Fully Paid	debt_consolidation	22.05	...	1

2 rows × 33 columns



In [77]:

```
Data.select_dtypes(['object']).columns
```

Out[77]:

```
Index(['sub_grade', 'verification_status', 'loan_status', 'purpose',  
      'initial_list_status', 'application_type'],  
      dtype='object')
```

In [78]:

```
Dummies=pd.get_dummies(Data[['verification_status', 'purpose', 'initial_list_status',  
                             'application_type', 'sub_grade']],drop_first=True)  
Data=Data.drop(['verification_status', 'purpose',  
               'initial_list_status', 'application_type', 'sub_grade'],axis=1)  
Data=pd.concat([Data,Dummies],axis=1)
```

In [79]:

```
Data=Data.drop('loan_status',axis=1)
```

In [80]:

```
Data.head(2)
```

Out[80]:

	loan_amnt	term	int_rate	installment	annual_inc	dti	open_acc	pub_rec	revol_bal	revol_util	...	sub_grade_F1	sub_grade_F2
0	10000.0	36	11.44	329.48	117000.0	26.24	16.0	0.0	36369.0	41.8	...	0	0
1	8000.0	36	11.99	265.68	65000.0	22.05	17.0	0.0	20131.0	53.3	...	0	0

2 rows × 79 columns

In [81]:

```
Data.to_csv(r'C:\Users\chumj\Downloads\Data.csv',index=False,header=True)
```

## WE ARE GOOD TO START IMPLEMENTING OUR MEACHINE LEARNING MODELS AND ANN.

In [82]:

```
#Setting X and y variables for our different models,that is the features(X),and label(y)  
X=Data.drop('repaid_loans',axis=1)  
y=Data['repaid_loans']
```

In [83]:

```
from sklearn.model_selection import train_test_split
```

In [84]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=101)
```

In [85]:

```
#lets begins with logistics regression model.  
from sklearn.linear_model import LogisticRegression
```

In [86]:

```
modell = LogisticRegression()
```

```
modell.fit(X_train,y_train)
```

Out[86]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                    warm_start=False)
```

In [87]:

```
predictions = modell.predict(X_test)
```

In [88]:

```
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
```

In [89]:

```
print(classification_report(y_test,predictions))
print(confusion_matrix(y_test,predictions))
print(accuracy_score(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.26	0.00	0.00	23363
1	0.80	1.00	0.89	95203
accuracy			0.80	118566
macro avg	0.53	0.50	0.45	118566
weighted avg	0.70	0.80	0.72	118566

```
[[  37 23326]
 [ 107 95096]]
0.8023632407266839
```

In [90]:

```
#using Decision tress and random forest.
from sklearn.tree import DecisionTreeClassifier
```

In [91]:

```
dtree = DecisionTreeClassifier()
```

In [92]:

```
dtree.fit(X_train,y_train)
```

Out[92]:

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                       max_depth=None, max_features=None, max_leaf_nodes=None,
                       min_impurity_decrease=0.0, min_impurity_split=None,
                       min_samples_leaf=1, min_samples_split=2,
                       min_weight_fraction_leaf=0.0, presort='deprecated',
                       random_state=None, splitter='best')
```

In [93]:

```
predictions1 = dtree.predict(X_test)
```

In [94]:

```
from sklearn.metrics import classification_report, confusion_matrix
```

In [95]:

```
print(classification_report(y_test,predictions1))
print(confusion_matrix(y_test,predictions1))
print(accuracy_score(y_test,predictions1))
```

	precision	recall	f1-score	support
0	0.56	0.59	0.58	23363
1	0.90	0.89	0.89	95203
accuracy			0.83	118566
macro avg	0.73	0.74	0.73	118566
weighted avg	0.83	0.83	0.83	118566

```
[[13773  9590]
 [10643 84560]]
0.829352428183459
```

In [97]:


```
cm3=confusion_matrix(y_test,predictions1)
```

In [98]:

```
class_names=[0,1] # name of classes
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)
# create heatmap
sns.heatmap(pd.DataFrame(cm3), annot=True, fmt='g')
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
```

Out[98]:

Text(0.5, 257.44, 'Predicted label')



In [99]:

```
#Given a customer below we you offer loan to the person or not
import random
random_ind=random.randint(0,len(Data))
```

In [100]:

```
new_person4=Data.drop('repaid_loans',axis=1).iloc[random_ind]
```

In [101]:

```
new_person4
```

Out[101]:

```
loan_amnt      20000.00
term           36.00
int_rate       11.48
installment    659.37
annual_inc     105000.00
...
sub_grade_G1   0.00
sub_grade_G2   0.00
sub_grade_G3   0.00
sub_grade_G4   0.00
sub_grade_G5   0.00
Name: 146375, Length: 78, dtype: float64
```

In [102]:

```
dtree.predict(new_person4.values.reshape(1,78))
```

Out[102]:

```
array([1], dtype=int64)
```

In [103]:

```
#check if this person end up paying the loan
Data.iloc[random_ind]['repaid_loans']
```

Out[103]:

```
1.0
```

In [ ]: